

Seshu_Miriyala_Final_Project

March 26, 2021

```
[1]: '''This program analyzes the coffee consumption, production, export and imports,
      among various countries and answers some key questions'''

# importing the required libraries
# Need Numpy to perform the array operations
import numpy as np
# We use pandas to read the csv files and load them into data frames
import pandas as pd
# we use matplotlib for plotting graphs and charts
import matplotlib as mp
import matplotlib.pyplot as pp
```

1 Load the data

```
[2]: # Loading coffee consumption data into data frame
# The data contains one column for the country name and one column each for the
      years 1990 till 2018
# The units of measure are in thousands 60-Kg bags
consumption_data = pd.read_csv('inputs/domestic-consumption.csv',delimiter=',')
# Merging any duplicates
consumption_data = consumption_data.groupby("domestic_consumption",
      as_index=True).sum()
# Checking the data after loadin into the data frame
consumption_data.head(5)
```

```
[2]:
```

	1990	1991	1992	1993	1994	\
domestic_consumption						
Angola	20.0	30.0	35.0	20.00	25.0	
Benin	0.0	0.0	0.0	0.00	0.0	
Bolivia (Plurinational State of)	25.0	27.0	27.5	28.50	29.5	
Brazil	8200.0	8500.0	8900.0	9100.00	9300.0	
Burundi	2.0	1.6	1.7	1.91	2.0	
	1995	1996	1997	1998	1999	\
domestic_consumption						
Angola	10.0	20.0	40.0	30.0	20.0	

Benin	0.0	0.0	0.0	0.0	0.0
Bolivia (Plurinational State of)	30.5	31.5	32.5	33.0	34.0
Brazil	10100.0	11000.0	11500.0	12200.0	12700.0
Burundi	2.0	2.0	2.0	2.0	2.0
...	2009	2010	2011	2012	\
domestic_consumption	...				
Angola	30.000	30.0	30.0	30.0	
Benin	0.000	0.0	0.0	0.0	
Bolivia (Plurinational State of)	46.000	47.5	49.0	50.5	
Brazil	18390.000	19132.0	19720.0	20330.0	
Burundi	1.399	2.0	2.0	2.0	
	2013	2014	2015	2016	2017 \
domestic_consumption					
Angola	30.0	30.0	30.0	30.0	30.0
Benin	0.0	0.0	0.0	0.0	0.0
Bolivia (Plurinational State of)	52.0	53.5	55.0	57.0	58.5
Brazil	20085.0	20333.0	20508.0	21225.0	21997.0
Burundi	2.0	2.0	2.0	2.0	2.0
	2018				
domestic_consumption					
Angola	30.0				
Benin	0.0				
Bolivia (Plurinational State of)	60.0				
Brazil	22250.0				
Burundi	2.0				

[5 rows x 29 columns]

```
[3]: # Loading the coffee production data into the data frames
# The data contains one column for the country name and one column each for the
# years 1990 till 2018
# The units of measure are in thousands 60-Kg bags
production_data = pd.read_csv('inputs/total-production.csv', delimiter=',')
production_data = production_data.groupby("total_production", as_index=True).
    sum()
# Checking the data after loading
production_data.head(5)
```

```
[3]:
total_production
Angola          50.3450    79.3310    77.5200
Benin           0.0000    0.0000    1.8050
Bolivia (Plurinational State of)  122.7770   103.5360   120.2350
Brazil         27285.6286  27293.4934  34603.3542
```

Burundi	487.3930	667.1990	620.2380	
	1993	1994	1995	\
total_production				
Angola	32.6080	76.802	62.1090	
Benin	0.0500	0.000	0.0000	
Bolivia (Plurinational State of)	50.8230	116.944	142.4850	
Brazil	28166.9786	28192.047	18060.2022	
Burundi	393.3540	664.143	433.9800	
	1996	1997	1998	\
total_production				
Angola	70.925	64.330	85.3440	
Benin	0.000	0.000	0.0000	
Bolivia (Plurinational State of)	124.579	140.719	137.9850	
Brazil	29196.743	26148.004	36760.8533	
Burundi	400.969	249.785	491.9920	
	1999	...	2009	2010 \
total_production		...		
Angola	54.9390	...	13.4200	34.9700
Benin	0.0000	...	0.0000	0.0000
Bolivia (Plurinational State of)	157.7020	...	128.4751	117.2249
Brazil	47577.8065	...	43976.8120	55428.4102
Burundi	350.5500	...	111.6130	352.9776
	2011	2012	2013	\
total_production				
Angola	28.7150	32.7900	34.9350	
Benin	0.0000	0.0000	0.0000	
Bolivia (Plurinational State of)	131.8354	105.2812	119.9122	
Brazil	48591.8289	55418.0012	54688.9664	
Burundi	204.1328	405.9615	163.2177	
	2014	2015	2016	\
total_production				
Angola	39.4050	40.5150	44.8300	
Benin	0.0000	0.0000	0.0000	
Bolivia (Plurinational State of)	99.8766	84.2191	77.9835	
Brazil	53304.7669	52870.5876	56788.1784	
Burundi	247.5500	274.1017	248.7933	
	2017	2018		
total_production				
Angola	35.0060	40.3874		
Benin	0.0000	0.0000		
Bolivia (Plurinational State of)	83.8112	82.5687		

Brazil	52739.8635	62924.8836
Burundi	202.1079	178.4206

[5 rows x 29 columns]

```
[4]: # Loading the prices paid to grower in each country into the data frame
# The data contains one column for the country name and one column each for the
# years 1990 till 2018
# The units of measure are in USD/Kg
prices_paid_to_farmers = pd.read_csv('inputs/prices-paid-to-growers.csv',
# delimiter=',')
prices_paid_to_farmers = prices_paid_to_farmers.
# groupby("prices_paid_to_growers", as_index=True).mean()
# Displaying the first 5 rows for verification
prices_paid_to_farmers.head(5)
```

```
[4]:
```

	1990	1991	1992	1993	1994	\
prices_paid_to_growers						
Brazil	0.984254	0.789049	0.803461	0.975950	2.207448	
Colombia	1.534724	1.481790	1.204656	1.106477	1.898327	
Dominican Republic	1.458168	1.382845	1.027841	1.172704	2.478234	
El Salvador	1.116194	0.983322	0.682322	0.780397	2.191177	
Ethiopia	1.348565	1.505322	1.351128	1.362442	2.418234	
	1995	1996	1997	1998	1999	...
prices_paid_to_growers						...
Brazil	2.351244	1.833606	2.464297	2.019828	1.445841	...
Colombia	2.199185	2.065245	2.939673	2.253433	1.906905	...
Dominican Republic	2.412000	1.935342	3.673587	2.102616	1.637609	...
El Salvador	2.218826	1.656073	2.762552	1.877384	1.256563	...
Ethiopia	2.539011	1.495287	2.010060	2.027901	1.449113	...
	2009	2010	2011	2012	2013	\
prices_paid_to_growers						
Brazil	1.922988	2.298156	3.637459	2.833251	2.068699	
Colombia	3.067442	3.985616	5.290868	3.679737	2.514539	
Dominican Republic	2.340731	3.071276	4.217812	3.262839	3.112084	
El Salvador	1.748205	2.425620	4.122256	2.651658	2.103276	
Ethiopia	1.701150	1.886530	3.211088	2.269088	1.618110	
	2014	2015	2016	2017	2018	
prices_paid_to_growers						
Brazil	2.382567	1.964829	2.214220	2.303390	1.746153	
Colombia	3.525662	2.635854	2.727530	2.776185	2.505914	
Dominican Republic	4.052172	3.775011	3.911395	3.826974	3.391053	
El Salvador	2.582929	1.930439	1.976442	1.934159	1.645300	
Ethiopia	1.837430	1.678360	1.927316	1.799278	1.560234	

[5 rows x 29 columns]

```
[5]: # Loading the retail price of coffee in different countries
# The data contains one column for the country name and one column each for the
↳ years 1990 till 2018
# The units of measure are in USD/Kg
coffee_retail_price = pd.read_csv('inputs/retail-prices.csv', delimiter=',')
coffee_retail_price = coffee_retail_price.groupby("retail_prices",
↳ as_index=True).mean()
# Displaying the first 5 rows for verification
coffee_retail_price.head(5)
```

```
[5]:
```

	1990	1991	1992	1993	1994	\
retail_prices						
Austria	10.816777	10.088300	11.015453	10.971302	10.110375	
Cyprus	6.247241	6.181015	6.335541	5.739514	7.019868	
Denmark	8.410596	8.101545	8.366446	7.682119	9.823400	
Finland	6.578366	6.004415	5.430464	4.282561	6.026490	
France	8.233996	7.571744	5.099338	4.481236	5.298013	
	1995	1996	1997	1998	1999	...
retail_prices						...
Austria	11.434879	11.964680	9.646799	8.763797	7.240618	...
Cyprus	9.403974	9.116998	8.918322	10.176600	9.690949	...
Denmark	12.295806	10.618102	10.949227	10.860927	8.675497	...
Finland	8.763797	7.108168	7.726269	7.549669	5.739514	...
France	8.145695	7.284768	6.092715	6.136865	5.629139	...
	2009	2010	2011	2012	2013	\
retail_prices						
Austria	15.342163	14.768212	18.366446	18.498896	19.028698	
Cyprus	12.207506	11.501104	13.377483	14.039735	14.282561	
Denmark	11.677704	12.008830	15.275938	14.834437	14.039735	
Finland	7.748344	8.145695	11.832230	10.596026	9.470199	
France	8.366446	7.991170	9.116998	8.653422	8.653422	
	2014	2015	2016	2017	2018	
retail_prices						
Austria	19.050773	16.423841	12.450331	13.730684	14.635762	
Cyprus	14.304636	11.699779	11.699779	12.141280	12.781457	
Denmark	14.194260	12.913907	10.905077	11.103753	11.699779	
Finland	9.514349	8.609272	8.101545	9.050773	9.359823	
France	8.322296	6.865342	7.196468	7.505519	8.123620	

[5 rows x 29 columns]

```
[6]: # Loading the exports of coffee crop by different countries
# The data contains one column for the country name and one column each for the
↳ years 1990 till 2018
# The units of measure are in thousands 60-Kg bags
coffee_exports = pd.read_csv('inputs/exports-crop-year.csv', delimiter=',')
coffee_exports = coffee_exports.groupby("exports_crop_year", as_index=True).
↳ sum()
# Displaying the first 5 rows for verification
coffee_exports.head(5)
```

```
[6]:
```

	1990	1991	1992 \
exports_crop_year			
Angola	79.3450	74.3310	67.5200
Benin	0.0000	0.0000	1.8050
Bolivia (Plurinational State of)	111.9770	82.9360	100.9350
Brazil	17862.6286	21808.4934	16752.3542
Burundi	412.3930	762.4910	671.6460

	1993	1994	1995 \
exports_crop_year			
Angola	27.6080	11.802	48.1090
Benin	0.0500	0.000	0.0000
Bolivia (Plurinational State of)	36.5230	85.944	110.4850
Brazil	18760.9786	15958.047	13760.2022
Burundi	352.8700	580.127	464.0700

	1996	1997	1998 \
exports_crop_year			
Angola	50.925	54.330	55.3440
Benin	0.000	0.000	0.0000
Bolivia (Plurinational State of)	105.079	109.219	100.9850
Brazil	17259.743	15352.004	21084.8533
Burundi	185.636	546.034	391.8500

	1999 ...	2009	2010 \
exports_crop_year	...		
Angola	39.9390 ...	3.4200	4.9700
Benin	0.0000 ...	0.0000	0.0000
Bolivia (Plurinational State of)	122.7020 ...	82.4751	69.7249
Brazil	21185.8065 ...	30254.8120	34054.4102
Burundi	437.7990 ...	172.9370	350.7196

	2011	2012	2013 \
exports_crop_year			
Angola	8.7150	7.7900	4.9350
Benin	0.0000	0.0000	0.0000
Bolivia (Plurinational State of)	82.8354	54.7812	67.9122

Brazil	32148.8289	29283.0012	32751.9664
Burundi	202.1328	405.9615	159.2177
	2014	2015	2016 \
exports_crop_year			
Angola	9.4050	10.5150	14.8300
Benin	0.0000	0.0000	0.0000
Bolivia (Plurinational State of)	46.3766	29.2191	20.9835
Brazil	37781.7669	37472.5876	33491.1784
Burundi	245.5500	274.1017	246.7933
	2017	2018	
exports_crop_year			
Angola	5.0060	10.3874	
Benin	0.0000	0.0000	
Bolivia (Plurinational State of)	25.3112	22.5687	
Brazil	30782.8635	37613.8836	
Burundi	195.1079	179.9206	

[5 rows x 29 columns]

1.1 Cleaning the data

```
[7]: consumption_data.isnull().sum().sum()
```

```
[7]: 0
```

```
[8]: production_data.isnull().sum().sum()
```

```
[8]: 0
```

```
[9]: prices_paid_to_farmers.isnull().sum().sum()
```

```
[9]: 0
```

```
[10]: coffee_retail_price.isnull().sum().sum()
```

```
[10]: 0
```

```
[11]: coffee_exports.isnull().sum().sum()
```

```
[11]: 0
```


1.2 Graph of brazil production data

In this section we try to graph the production of Brazil over the years

```
[12]: # Defining a function to draw line chart using matplotlib so that we can reuse
      ↪ the code
def plotLineGraph(df, countriesList, colorList):
    x = np.arange(len(df.columns))
    for country in countriesList:
        t = np.array(df[df.index == country])[0]
        pp.plot(x, t, label=country, color=colorList[countriesList.
      ↪ index(country)])
    pp.xticks(x,np.array(df.columns))
    pp.xticks(rotation=90)
    pp.legend()
    pp.show()

# Default colors used in the code
colorsList = list(['#1f77b4', 'red', 'green'])

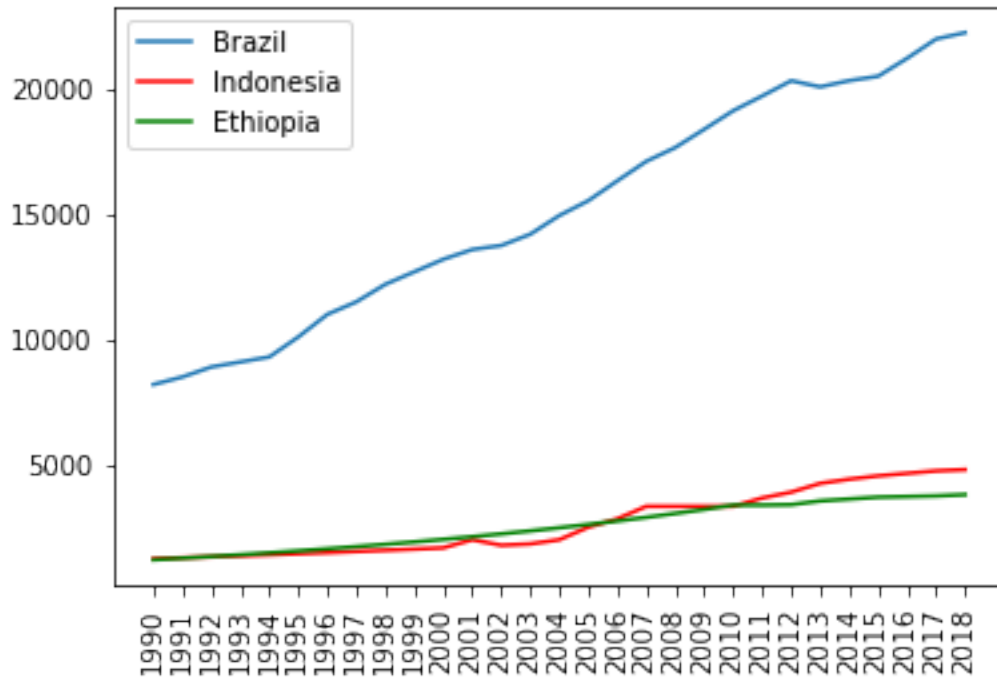
[13]: # Let us now see how are the production numbers for the coffee crop they grew
      ↪ in Brazil, Columbia and Viet Nam
plotLineGraph(production_data, list(['Brazil', 'Colombia', 'Viet Nam']),
      ↪ colorsList)
```



1.3 Graph of brazil consumption data

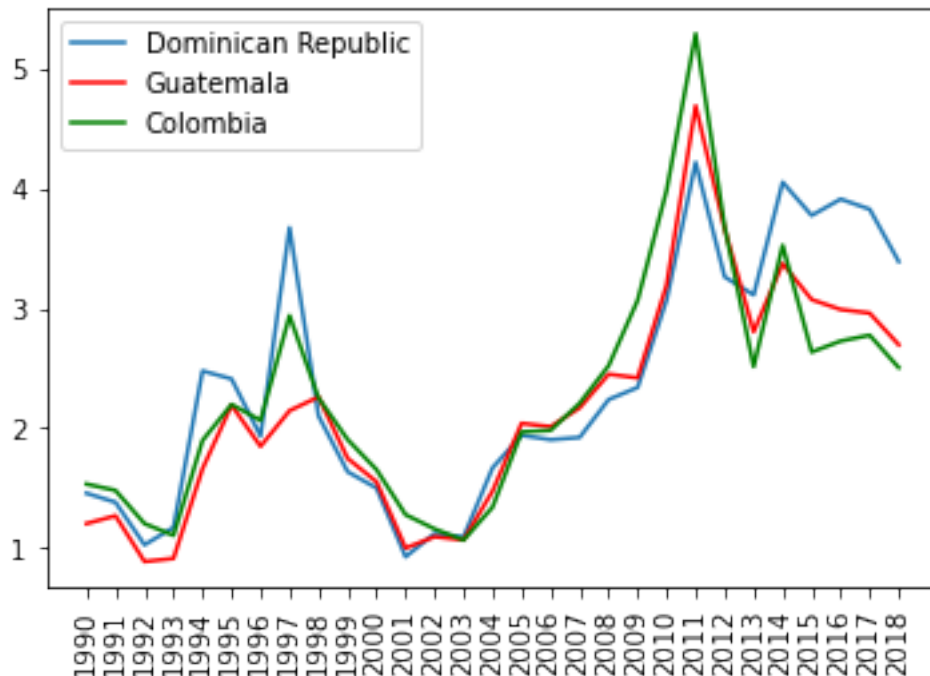
In this section we try to graph the consumption of Brazil over the years


```
[14]: # Let us now see how are the production numbers for the coffee crop they grew
      ↪ in Brazil, Columbia and Viet Nam
      plotLineGraph(consumption_data, list(['Brazil', 'Indonesia', 'Ethiopia']),
      ↪ colorsList)
```



As we can see, Consumption and production are high in Brazil than any other country

```
[15]: # Let us now see how are the farmers paid for the coffee crop they grew in
      ↪ Brazil, Columbia and India
      plotLineGraph(prices_paid_to_farmers, list(['Dominican Republic', 'Guatemala',
      ↪ 'Colombia']), list(['#1f77b4', 'red', 'green']))
```



As we can see Prices of the top coffee expensive countries varies by time. In the recent years prices of coffee are trading higher in Dominican Republic than other two countries

1.4 Taking average of last five years

In order to better analyze the data and to eliminate noise I have chosen to take average of last five years of data

```
[16]: # Defining a function to compute Five Year Average
def computeFiveYearAvg(df):
    # We then sum up the values of the last five years and divide the sum by 5
    →to get the average
    df['Five_year_avg'] = (df['2014'] + df['2015'] + df['2016'] + df['2017'] +
    →df['2018'])/5

    # consumption_data.reset_index(inplace=True)
    # We then create a new data frame with the calculated values
    avg_df = pd.DataFrame(data={'Country': np.array(df.index), 'Five_year_avg':
    →np.array(df.Five_year_avg)})
    avg_df.set_index('Country', inplace=True)

    return avg_df
```

```
[17]: # Computing the Five Year avg for consumption data
avg_consumption_data = computeFiveYearAvg(consumption_data)
```

```
# Checking if the data frame is right
avg_consumption_data.head(5)
```

```
[17]:
```

	Five_year_avg
Country	
Angola	30.0
Benin	0.0
Bolivia (Plurinational State of)	56.8
Brazil	21262.6
Burundi	2.0

```
[18]: # Computing the Five Year avg for production data
avg_production_data = computeFiveYearAvg(production_data)
# Checking if the data frame is right
avg_production_data.head(5)
```

```
[18]:
```

	Five_year_avg
Country	
Angola	40.02868
Benin	0.00000
Bolivia (Plurinational State of)	85.69182
Brazil	55725.65600
Burundi	230.19470

```
[19]: # Computing the Five Year avg for Prices paid to farmers data
avg_prices_paid_to_farmers = computeFiveYearAvg(prices_paid_to_farmers)
# Checking if the data frame is right
avg_prices_paid_to_farmers.head(5)
```

```
[19]:
```

	Five_year_avg
Country	
Brazil	2.122232
Colombia	2.834229
Dominican Republic	3.791321
El Salvador	2.013854
Ethiopia	1.760524

```
[20]: # Computing the Five Year avg for Coffee retail price data
avg_coffee_retail_price = computeFiveYearAvg(coffee_retail_price)
# Checking if the data frame is right
avg_coffee_retail_price.head(5)
```

```
[20]:
```

	Five_year_avg
Country	
Austria	15.258278
Cyprus	12.525386
Denmark	12.163355

Finland	8.927152
France	7.602649

```
[21]: # Computing the Five Year avg for coffee exports data
avg_coffee_exports = computeFiveYearAvg(coffee_exports)
# Checking if the data frame is right
avg_coffee_exports.head(5)
```

```
[21]:
```

Country	Five_year_avg
Angola	10.02868
Benin	0.00000
Bolivia (Plurinational State of)	28.89182
Brazil	35428.45600
Burundi	228.29470

1.5 Which country consumes more coffee?

For this we need to sort the consumption data and get the country that consumes coffee most

```
[22]: # Extracting the country which has highest coffee consumption
max_consumption_country = avg_consumption_data[avg_consumption_data.
↪Five_year_avg == np.max(avg_consumption_data.Five_year_avg)]
# Displaying the result
max_consumption_country
```

```
[22]:
```

Country	Five_year_avg
Brazil	21262.6

1.6 Which countries have excess coffee production or highest coffee deficit?

For this we need to take off the consumption numbers from the production for each country

```
[23]: # Extracting the Five Year Avg coffee production and consumption data
avg_consumption_data = avg_consumption_data.rename(columns={'Five_year_avg':
↪'Five_year_avg_consumption'})
avg_production_data=avg_production_data.rename(columns={'Five_year_avg':
↪'Five_year_avg_production'})

# Merging the two datasets using the Country as the join column
production_consumption = pd.merge(avg_consumption_data,avg_production_data,
↪on='Country')
# Displaying the results of the join
production_consumption.head(5)
```

```
[23]:
```

	Five_year_avg_consumption \
Country	
Angola	30.0
Benin	0.0
Bolivia (Plurinational State of)	56.8
Brazil	21262.6
Burundi	2.0

	Five_year_avg_production
Country	
Angola	40.02868
Benin	0.00000
Bolivia (Plurinational State of)	85.69182
Brazil	55725.65600
Burundi	230.19470

```
[24]: # Subtracting the consumption values from production values for each country to
      ↪ get the excess coffee the countries are producing
production_consumption['excess_production'] = production_consumption.
      ↪ Five_year_avg_production - production_consumption.Five_year_avg_consumption

# Calculating the mean of production and consumption values across countries
production_consumption['mean_production'] = np.mean(production_consumption.
      ↪ Five_year_avg_production)
production_consumption['mean_consumption'] = np.mean(production_consumption.
      ↪ Five_year_avg_consumption)

# Filtering the countries which are deficit in coffee production
Excess_producing_countries = production_consumption[production_consumption.
      ↪ excess_production > 0]

# Displaying the excess coffee production in tabular form
production_consumption.head(5)
```

```
[24]:
```


	Five_year_avg_consumption \
Country	
Angola	30.0
Benin	0.0
Bolivia (Plurinational State of)	56.8
Brazil	21262.6
Burundi	2.0

	Five_year_avg_production	excess_production \
Country		
Angola	40.02868	10.02868
Benin	0.00000	0.00000
Bolivia (Plurinational State of)	85.69182	28.89182

Brazil	55725.65600	34463.05600
Burundi	230.19470	228.19470

	mean_production	mean_consumption
Country		
Angola	2830.954115	867.873341
Benin	2830.954115	867.873341
Bolivia (Plurinational State of)	2830.954115	867.873341
Brazil	2830.954115	867.873341
Burundi	2830.954115	867.873341

```
[25]: pp.bar(x=np.array(production_consumption.index), height=np.
      ↪array(production_consumption.excess_production), align='center',
      ↪label='Excess coffee production')
pp.xticks(rotation = 90)
pp.plot(np.array(production_consumption.index), production_consumption.
      ↪mean_production, color='orange', label="Avg production")
pp.plot(np.array(production_consumption.index), production_consumption.
      ↪mean_consumption, color='red', label="Avg Consumption")
pp.legend()
pp.show()
```



```
[26]: # Extracting the country that has maximum coffee surplus production
max_surplus_country = production_consumption[production_consumption.
    ↳excess_production == np.max(production_consumption.excess_production)]

# Displaying the result
max_surplus_country
```

```
[26]:      Five_year_avg_consumption  Five_year_avg_production  \
Country
Brazil                21262.6                55725.656

      excess_production  mean_production  mean_consumption
Country
Brazil           34463.056           2830.954115           867.873341
```


```
[27]: # Extracting the country that has maximum coffe deficit
max_deficit_country = production_consumption[production_consumption.
↳excess_production == np.min(production_consumption.excess_production)]
# Displaying the results
max_deficit_country
```

```
[27]:          Five_year_avg_consumption  Five_year_avg_production \
Country
Philippines                3062.0                205.09132

          excess_production  mean_production  mean_consumption
Country
Philippines        -2856.90868        2830.954115        867.873341
```

1.7 Which country has best coffee production

```
[28]: production_data['mean_production_2018'] = np.mean(production_data['2018'])
pp.bar(np.array(production_data.index), np.array(production_data['2018']))
pp.xticks(rotation=90)
pp.plot(np.array(production_data.index), np.array(production_data.
↳mean_production_2018), color='orange', label='Avg Production')
pp.legend()
pp.show()
```

```
[29]: # Extracting the country having the highest coffee production
max_production_country = production_data[production_data['2018'] == np.
    ↳max(production_data['2018'])]
# Displaying the results
max_production_country
```

```
[29]:
```

	1990	1991	1992	1993	1994 \
total_production					
Brazil	27285.6286	27293.4934	34603.3542	28166.9786	28192.047

	1995	1996	1997	1998	1999 \
total_production					
Brazil	18060.2022	29196.743	26148.004	36760.8533	47577.8065

	...	2011	2012	2013	2014	\
total_production	...					
Brazil	...	48591.8289	55418.0012	54688.9664	53304.7669	

		2015	2016	2017	2018	\
total_production						
Brazil		52870.5876	56788.1784	52739.8635	62924.8836	

		Five_year_avg	mean_production_2018
total_production			
Brazil		55725.656	3052.441745

[1 rows x 31 columns]

```
[30]: # Extracting the country which has minimum coffee production
min_production_country = production_data[production_data['2018'] == np.
    ↳min(production_data['2018'])]
# Displaying the results
min_production_country
```

```
[30]:
```

		1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	\
total_production												
Benin		0.0	0.0	1.805	0.05	0.0	0.0	0.0	0.0	0.0	0.0	

	...	2011	2012	2013	2014	2015	2016	2017	2018	\
total_production	...									
Benin	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

		Five_year_avg	mean_production_2018
total_production			
Benin		0.0	3052.441745

[1 rows x 31 columns]

It seems Brazil has the highest coffee production which is 62925 thousand 60Kg bags while Benin has not produced any coffee since 1990

1.8 How are farmers paid for the coffee they grow in each country

For this lets compare the crop price, retail price of coffee across countries

```
[31]: # Extracting the prices paid to farmers in 2018 and using the logarithmic
    ↳function to reduce the gap
prices_paid = np.log(prices_paid_to_farmers[['2018']])
# Extracting the production values for 2018
coffee_production = np.log(production_data[['2018']])

# Renaming the column names to be clear during the merge operation
```

```

prices_paid = prices_paid.rename(columns={'2018': 'price_paid_to_growers_2018'})
coffee_production = coffee_production.rename(columns={'2018': 'production_2018'})

#Joining the two datasets on Country
production_prices = prices_paid.join(coffee_production, how='inner')

# Displaying the result
production_prices.head(5)

```

```

[31]:

```

	price_paid_to_growers_2018	production_2018
Brazil	0.557415	11.049697
Colombia	0.918654	9.536615
Dominican Republic	1.221140	6.065899
El Salvador	0.497923	6.634532
Ethiopia	0.444836	8.958826

```

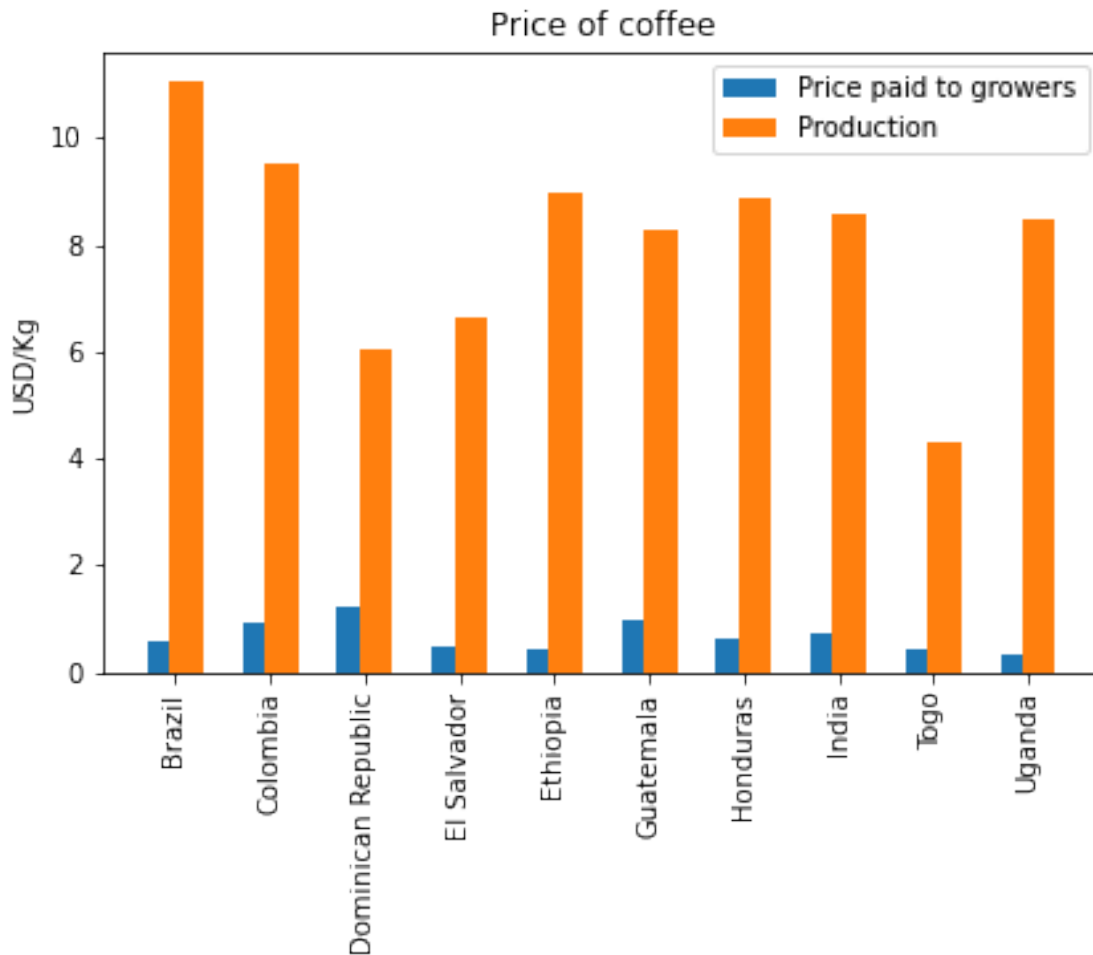
[32]: x = np.arange(len(production_prices.index)) # the label locations
width = 0.35 # the width of the bars

fig, ax = pp.subplots()
rects1 = ax.bar(x - width/3, production_prices.price_paid_to_growers_2018, width,
    label='Price paid to growers')
rects2 = ax.bar(x + width/3, production_prices.production_2018, width,
    label='Production')

# Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_ylabel('USD/Kg')
ax.set_title('Price of coffee')
ax.set_xticks(x)
ax.set_xticklabels(production_prices.index)
ax.legend()

fig.tight_layout()
pp.xticks(rotation=90)
pp.legend()
pp.show()

```



1.9 Which countries export most of the coffee?

```
[33]: avg_coffee_exports.head(5)
```

```
[33]:
```

Country	Five_year_avg
Angola	10.02868
Benin	0.00000
Bolivia (Plurinational State of)	28.89182
Brazil	35428.45600
Burundi	228.29470

```
[34]: # Charting the Export values
avg_coffee_exports['mean_export_2018'] = np.mean(avg_coffee_exports.
↪Five_year_avg)
```

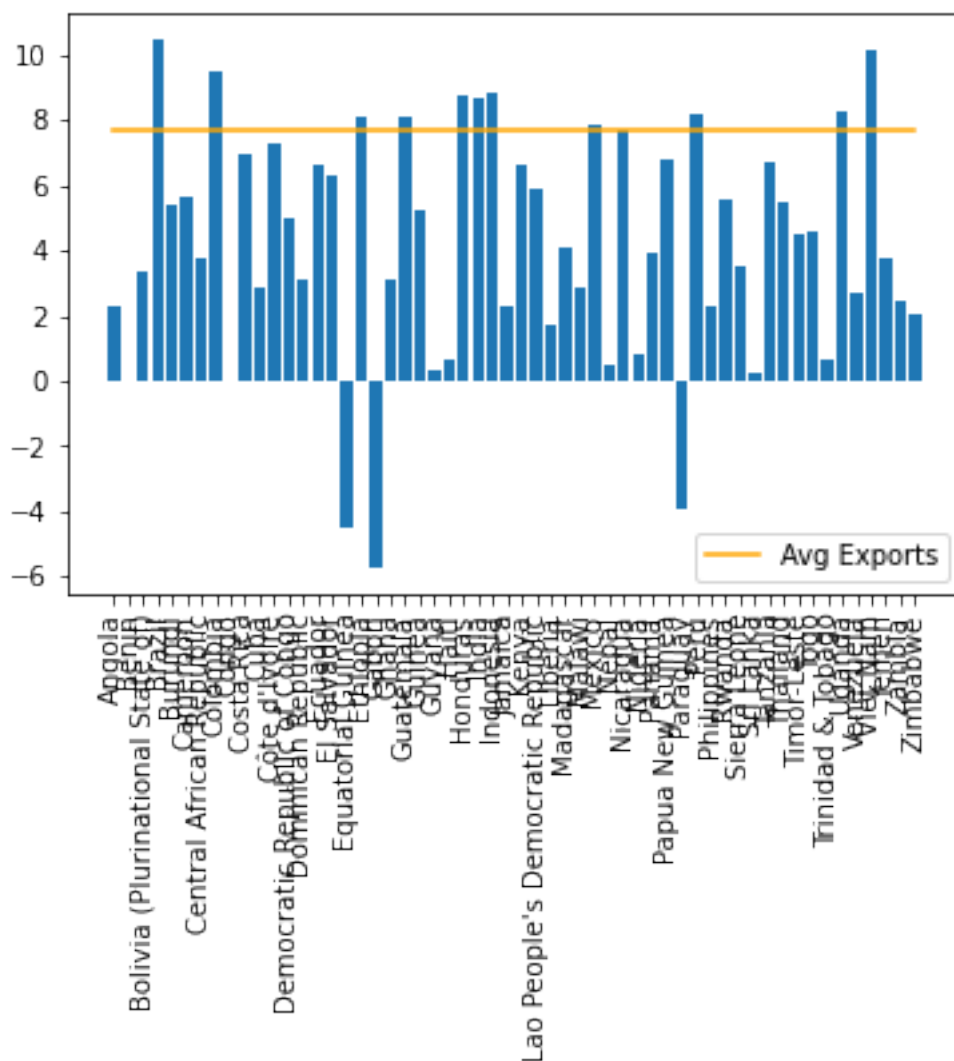
```

pp.bar(np.array(avg_coffee_exports.index), np.array(np.log(avg_coffee_exports.
↪Five_year_avg)))
pp.xticks(rotation=90)
pp.plot(np.array(avg_coffee_exports.index), np.array(np.log(avg_coffee_exports.
↪mean_export_2018)), color='orange', label='Avg Exports')
pp.legend()
pp.show()

```

/Users/seshumiriyala/opt/anaconda3/lib/python3.8/site-packages/pandas/core/arraylike.py:358: RuntimeWarning: divide by zero encountered in log

```
result = getattr(ufunc, method)(*inputs, **kwargs)
```



```
[35]: # Extracting the maximum coffee exporting country
max_exporting_country = avg_coffee_exports[avg_coffee_exports['Five_year_avg']_
↳== np.max(avg_coffee_exports['Five_year_avg'])]
# Displaying the results
max_exporting_country
```

```
[35]:          Five_year_avg  mean_export_2018
Country
Brazil          35428.456          2136.673436
```

1.10 Forecasting the 2019 production and consumption for Brazil

```
[36]: # Forecating the production and consumption values using simple moving average
consumption_data_transposed = consumption_data.T
consumption_data_transposed['SMA_3'] = consumption_data_transposed.iloc[:,3].
↳rolling(window=3).mean()
consumption_data_transposed.head(5)
```

```
[36]: domestic_consumption  Angola  Benin  Bolivia (Plurinational State of)  Brazil  \
1990                    20.0    0.0                                25.0  8200.0
1991                    30.0    0.0                                27.0  8500.0
1992                    35.0    0.0                                27.5  8900.0
1993                    20.0    0.0                                28.5  9100.0
1994                    25.0    0.0                                29.5  9300.0
```

```
domestic_consumption  Burundi  Cameroon  Central African Republic  Colombia  \
1990                    2.00    83.300                                28.0  1235.0
1991                    1.60    83.333                                28.0  1269.0
1992                    1.70    83.333                                28.0  1303.0
1993                    1.91   100.000                                25.0  1339.0
1994                    2.00   100.000                                15.0  1375.0
```

```
domestic_consumption  Congo  Costa Rica  ...  Timor-Leste  Togo  \
1990                3.000    375.0  ...          0.0    1.0
1991                2.667    375.0  ...          0.0    1.0
1992                2.667    375.0  ...          0.0    1.0
1993                3.000    375.0  ...          0.0    1.0
1994                3.000    375.0  ...          0.0    1.0
```


```
domestic_consumption  Trinidad & Tobago  Uganda  Venezuela  Viet Nam  Yemen  \
1990                                8.0    70.0    782.39    150.0    0.0
1991                               11.5    75.0    815.48    230.0    0.0
1992                               10.0    75.0    849.97    250.0    0.0
1993                               10.0    75.0    885.93    267.0    0.0
1994                               14.0    80.0    923.40    267.0    0.0
```

```
domestic_consumption  Zambia  Zimbabwe          SMA_3
```

1990	1.500	8.000	NaN
1991	1.500	8.000	NaN
1992	1.432	8.000	8533.333333
1993	1.000	8.333	8833.333333
1994	1.500	8.333	9100.000000

[5 rows x 57 columns]

```
[37]: # Let us now use the data to predict the coffee consumption for Brazil in 2019
      ↪ and chart it
consumption_data_normal = consumption_data_transposed.T
del consumption_data_normal['Five_year_avg']
plotLineGraph(consumption_data_normal, list(['Brazil', 'SMA_3']),
      ↪ list(['#1f77b4', 'red']))
```



```
[38]: # Consumption values for Brazil in 2019 is
brazil_consumption_2019 = (consumption_data_normal[consumption_data_normal.
      ↪ index == 'SMA_3']['2018'] + consumption_data_normal[consumption_data_normal.
      ↪ index == 'SMA_3']['2017'] + consumption_data_normal[consumption_data_normal.
      ↪ index == 'SMA_3']['2016'])/3
```

```
[39]: # Let us now see how are the farmers paid for the coffee crop they grew in
      ↪ Brazil, Columbia and India
# Forecating the production values using simple moving average
```


```

production_data_transposed = production_data.T
production_data_transposed['SMA_3'] = production_data_transposed.iloc[:,3].
    ↪rolling(window=3).mean()

production_data_normal = production_data_transposed.T

del production_data_normal['Five_year_avg']
del production_data_normal['mean_production_2018']
plotLineGraph(production_data_normal, list(['Brazil', 'SMA_3']),
    ↪list(['#1f77b4', 'red']))

```



[40]: *# Production values for Brazil in 2019 is*

```

brazil_production_2019 = (production_data_normal[production_data_normal.index
    ↪== 'SMA_3']['2018'] + production_data_normal[production_data_normal.index ==
    ↪'SMA_3']['2017'] + production_data_normal[production_data_normal.index ==
    ↪'SMA_3']['2016'])/3

```

```

[41]: print("Coffee production in Brazil for 2019 is", brazil_production_2019.
    ↪values[0], "(In thousand 60-kg bags)")
print("Coffee consumption in Brazil for 2019 is", brazil_consumption_2019.
    ↪values[0], "(In thousand 60-kg bags)")

```

Coffee production in Brazil for 2019 is 55312.787544444444 (In thousand 60-kg bags)

Coffee consumption in Brazil for 2019 is 21252.0 (In thousand 60-kg bags)